

COMP7630 – Web Intelligence and its Applications

Singular Value Decomposition

Valentino Santucci

(valentino.santucci@unistrapg.it)

Singular Value Decomposition (SVD)

- Eigendecomposition does not work with rectangular matrices (and also with non-diagonalizable matrices)
- SVD generalizes EigenDecomposition on rectangular matrices (to be precise: on any kind of matrix)
- It is also possible to describe PCA by using SVD, though we will not do it.
 - The trick is to work directly with data-table X and not with $X^T X$ as seen before.
 - ScikitLearn uses the SVD implementation, though it is mathematically equivalent to what we have seen!

Singular Value Decomposition (SVD)

- Another way to factorize a matrix into **singular vectors** and **single values**.
- Every real matrix (**even not a square**) has a SVD.
- Singular value decomposition

$$\mathbf{A} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T$$

- $\mathbf{A} \in \mathbb{R}^{m \times n}$ $\mathbf{U} \in \mathbb{R}^{m \times m}$ $\mathbf{\Sigma} \in \mathbb{R}^{m \times n}$ $\mathbf{V} \in \mathbb{R}^{n \times n}$
- \mathbf{U} and \mathbf{V} are orthogonal matrices.
- $\mathbf{\Sigma}$ is diagonal (but not necessarily square) and the elements along the diagonal are the **singular values**.
- \mathbf{U} is **left-singular vector** and \mathbf{V} is **right-singular vector**.

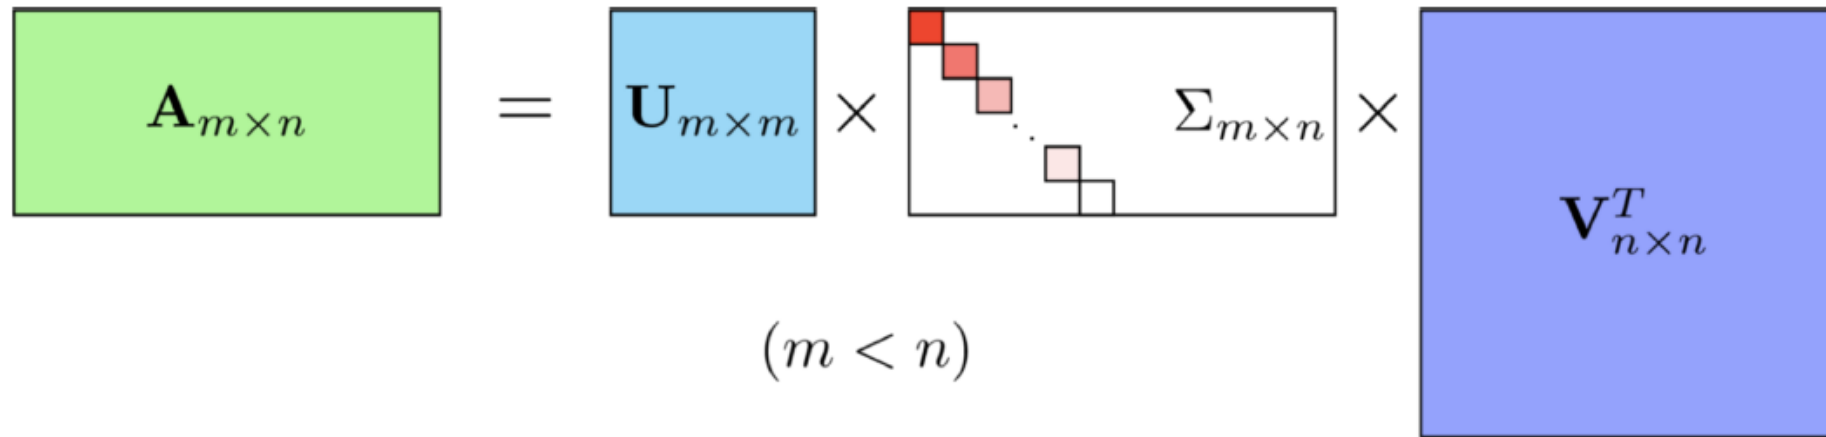
Some properties of the SVD

- U columns are the eigenvectors of $A^T A$ (which is symmetric by definition)
- They are called "left singular vectors" of A
- V^T rows are the eigenvectors of AA^T
- They are called "right singular vectors" of A
- The non-zero diagonal values in Σ are the square-root of the eigenvalues of both $A^T A$ and AA^T (which are equal)
- They are called "singular values" of A

Singular Value Decomposition (SVD)

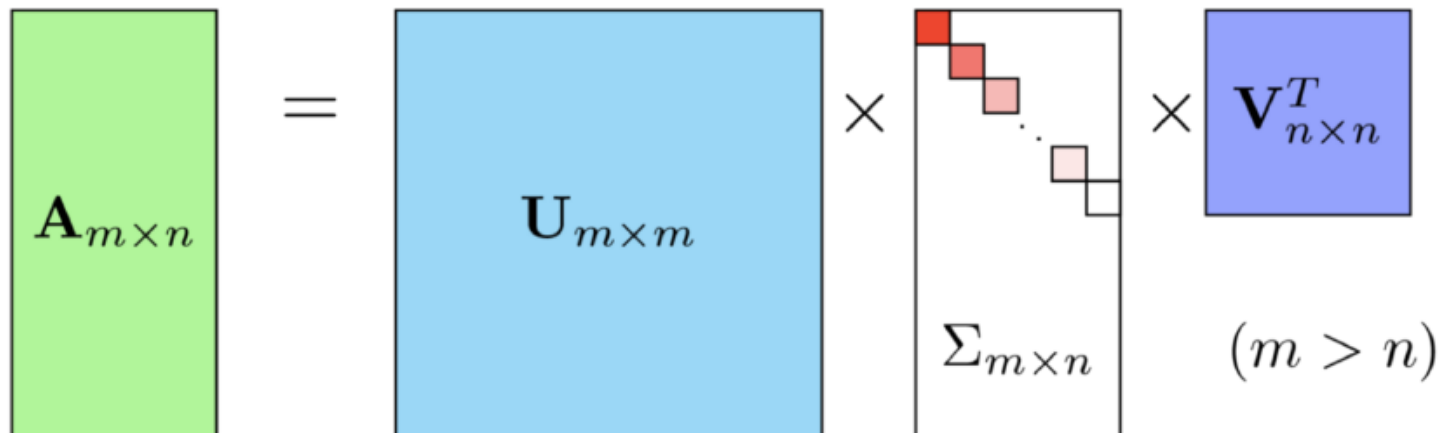
$$\mathbf{A}_{m \times n} = \mathbf{U}_{m \times m} \times \Sigma_{m \times n} \times \mathbf{V}_{n \times n}^T$$

$(m < n)$


The diagram illustrates the SVD decomposition for a matrix A of size m x n where m is less than n. Matrix A is shown as a green rectangle. It is equal to the product of three matrices: U (a blue square of size m x m), Sigma (a white rectangle of size m x n with a diagonal of red and pink squares), and V^T (a blue square of size n x n).

$$\mathbf{A}_{m \times n} = \mathbf{U}_{m \times m} \times \Sigma_{m \times n} \times \mathbf{V}_{n \times n}^T$$

$(m > n)$

The diagram illustrates the SVD decomposition for a matrix A of size m x n where m is greater than n. Matrix A is shown as a green rectangle. It is equal to the product of three matrices: U (a blue square of size m x m), Sigma (a white rectangle of size m x n with a diagonal of red and pink squares), and V^T (a blue square of size n x n).

Truncated SVD



$$\mathbf{A} \approx \tilde{\mathbf{A}}_k = \mathbf{U}_k \Sigma_k \mathbf{V}_k^T$$

Truncated SVD

- Keep only the k largest singular values.

$$\mathbf{A} \approx \tilde{\mathbf{A}}_k = \mathbf{U}_k \mathbf{\Sigma}_k \mathbf{V}_k^T$$

- $\tilde{\mathbf{A}}_k \in \mathbb{R}^{m \times n}$ $\mathbf{U}_k \in \mathbb{R}^{m \times k}$ $\mathbf{\Sigma}_k \in \mathbb{R}^{k \times k}$ $\mathbf{V}_k \in \mathbb{R}^{k \times n}$

- It can be shown that this gives the minimum value for the Frobenius norm of $\|\mathbf{A} - \tilde{\mathbf{A}}_k\|_F$

Truncated SVD

- It is important to note that:
 - We reduce the number of columns in the "left-matrix" U
 - We reduce the number of rows in the "right-matrix" V
 - But \tilde{A} has the same shape of A
- Anyway, \tilde{A} has a smaller rank than A
(recall: the rank of a matrix is the maximum number of columns which are linearly independent to each other)
- Why this may be useful?
 - A can be interpreted as a data-matrix containing noise
 - \tilde{A} is a denoised version of A
 - U_k rows are denoised/reduced representations of A rows
 - V_k^T columns are denoised/reduced representations of A columns

Another perspective on SVD

$$U\Sigma V^T x = U\Sigma \begin{bmatrix} v_1^T x \\ \vdots \\ v_m^T x \end{bmatrix} = U \begin{bmatrix} \sigma_1 v_1^T x \\ \vdots \\ \sigma_m v_m^T x \end{bmatrix} = \sum_k u_k \sigma_k v_k^T x = \sum_k \sigma_k u_k v_k^T x.$$

Hence $U\Sigma V^T = \sum_k \sigma_k u_k v_k^T$.

- The **SVD decomposition** can be rewritten as a sum of rank-1 matrices: those obtained by the outer product between the k -th column of U and the k -th row of V^T , weighted by k -th singular value.
 - Recall: the outer product of two vectors returns a rank-1 matrix) and, obviously, multiplying by a scalar does not change the rank of a matrix.
- So, the **Truncated SVD** acts as removing the terms with smaller weights in the summation!!!
 - (If you are familiar with it, it is a sort of "discrete Fourier transform")

SVD in Python

```
In [7]: import numpy as np

In [8]: from scipy.linalg import svd

In [9]: X = np.array([ [1,2,1,2],
...:                   [0,1,0,1],
...:                   [1,0,1,0],
...:                   [1,2,3,4] ])

In [10]: #SVD:  $X = U * S * V^T$ 

In [11]: U, s, VT = svd(X)

In [12]: U
Out[12]:
array([[ -0.47547615, -0.35956946,  0.69135054,  0.40824829],
       [ -0.18137369, -0.54296386,  0.0750144 , -0.81649658],
       [ -0.11272877,  0.72635827,  0.54132174, -0.40824829],
       [ -0.85341563,  0.21978106, -0.47262887,  0.          ]])

In [13]: s
Out[13]: array([6.38105353e+00, 1.49005261e+00, 1.03048489e+00, 3.21362686e-16])

In [14]: VT
Out[14]:
array([[ -0.22592203, -0.4449355 , -0.49340627, -0.71241974],
       [  0.39365716, -0.55202122,  0.68865488, -0.2570235 ],
       [  0.737559   ,  0.49729767, -0.17973513, -0.41999646],
       [ -0.5        ,  0.5        ,  0.5        , -0.5        ]])

In [15]: X_recovered = U.dot(np.diag(s)).dot(VT)

In [16]: X_recovered
Out[16]:
array([[ 1.00000000e+00,  2.00000000e+00,  1.00000000e+00,
         2.00000000e+00],
       [-3.48525357e-17,  1.00000000e+00, -2.02744139e-16,
         1.00000000e+00],
       [ 1.00000000e+00, -1.87190742e-15,  1.00000000e+00,
        -1.48282707e-16],
       [ 1.00000000e+00,  2.00000000e+00,  3.00000000e+00,
         4.00000000e+00]])
```

Truncated SVD in Python

```
In [21]: for k in [1,2,3]:
...:     X_recovered = U[:, :k].dot(np.diag(s[:k])).dot(VT[:k, :])
...:     norm = np.linalg.norm(X-X_recovered)
...:     print('---')
...:     print(f'X{k} = \n{X_recovered}')
...:     print(f'norm(X,X{k}) = {norm}')
...:
---
X1 =
[[0.69  1.35  1.5   2.16]
 [0.26  0.51  0.57  0.82]
 [0.16  0.32  0.35  0.51]
 [1.23  2.42  2.69  3.88]]
norm(X,X1) = 1.811672121361281
---
X2 =
[[ 0.47  1.65  1.13  2.3 ]
 [-0.06  0.96  0.01  1.03]
 [ 0.59 -0.28  1.1   0.23]
 [ 1.36  2.24  2.91  3.8 ]]
norm(X,X2) = 1.0304848877206008
---
X3 =
[[ 1.00e+00  2.00e+00  1.00e+00  2.00e+00]
 [-1.66e-16  1.00e+00 -7.15e-17  1.00e+00]
 [ 1.00e+00 -1.81e-15  1.00e+00 -2.14e-16]
 [ 1.00e+00  2.00e+00  3.00e+00  4.00e+00]]
norm(X,X3) = 3.670575228292439e-15
```