

COMP7630 – Web Intelligence and its Applications

# Web Information Retrieval

## (part 1)

Valentino Santucci

([valentino.santucci@unistrapg.it](mailto:valentino.santucci@unistrapg.it))

# Outline

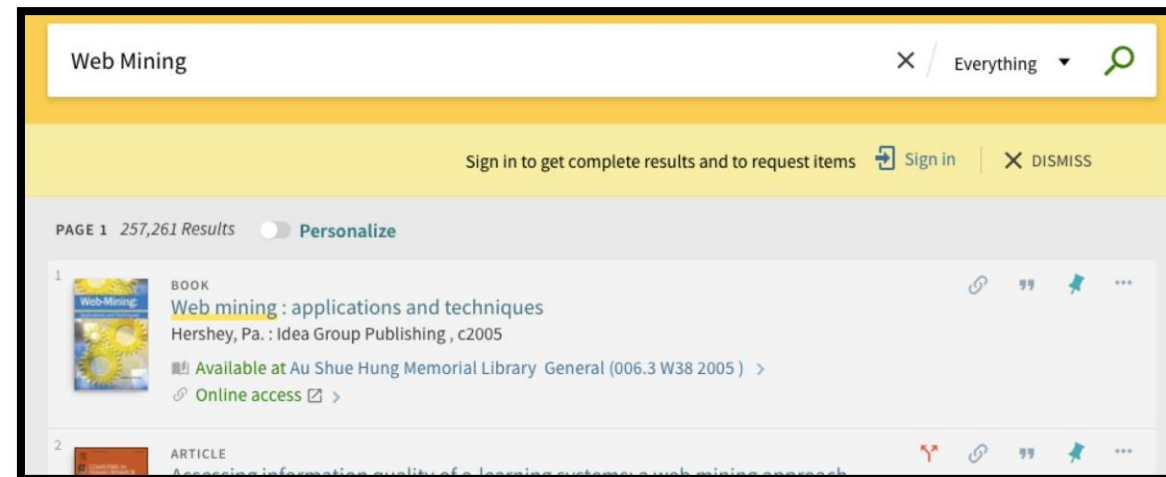
- Goals and architecture of a IR system
- IR models
- Word2Vec and other modern Language Models

# Information Retrieval vs Database Search

- Querying a relational database in a DBMS

```
1 select
2   f.contextid,
3   x.instanceid,
4   c.fullname as course_full_name,
5   c.shortname as course_short_name,
6   sum(f.filesize)/1000000000 as size_in_gigabytes,
7   sum(case when (f.filesize > 0) then 1 else 0 end) as number_of_files,
8   from_unixtime(c.timemodified) as last_update
9 from
10  mdl_files f inner join mdl_context x
11  on f.contextid = x.id
12  and x.contextlevel = 50
13  inner join mdl_course c
14  on c.id = x.instanceid
15 group by
16  f.contextid, x.instanceid
17 order by
18  sum(filesize) desc
```

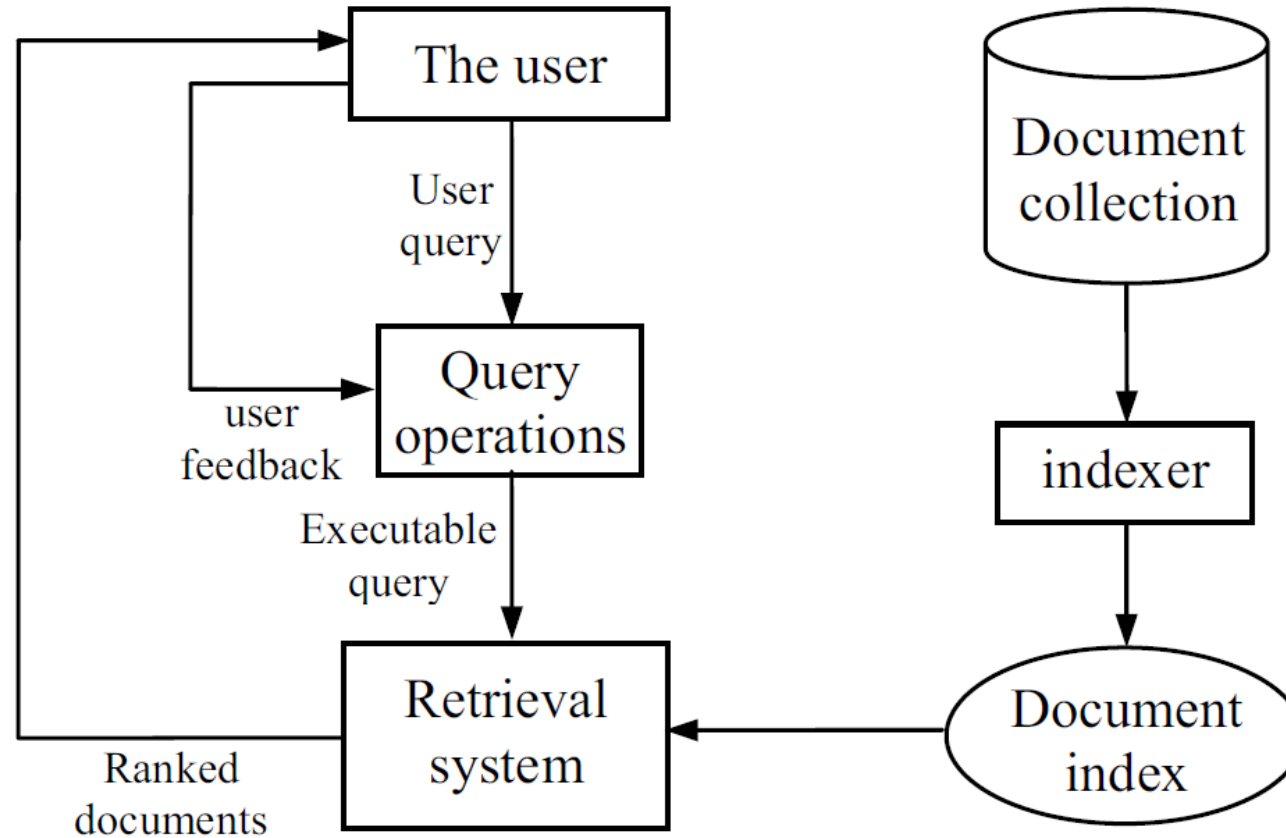
- Querying an IR system



# Information Retrieval vs Database Search

- Database Search (≠ Information Retrieval)
  - Basic information unit = a **data record**
  - Highly **structured data** and stored in relational tables
  - Results are ranked by **fields' values**
  - Common query format = **SQL statement**
- What is information retrieval?
  - Basic information unit = a **document (unstructured)**
  - A large collection of documents = the text “database” (or **corpus**).
  - Retrieval = finding a set of documents relevant to the **user query**.
  - The documents are ranked based on **relevance scores** to the query.
  - Common query format = a list of keywords (also called **terms**).

# The architecture of an IR System



# Forms of user queries

1. **Keyword queries**: list of terms (connected by a "soft" AND). It is not strictly necessary that retrieved document contains all the terms.
2. **Boolean queries**: terms and Boolean operators (AND, OR, NOT). Only exact matches.
3. **Phrase queries**: sequence of terms (ex: "web mining techniques")
4. **Proximity queries**: relaxed phrase queries where documents with query terms within close proximity have higher relevance.
5. **Full document queries**: the query is a document (ex: similar pages)
6. **Natural language questions**: the most complex form of query

# Components of IR systems

- Query operations module

In the simplest case, it does nothing but just pass the query to the retrieval engine after some **simple pre-processing**, e.g., removal of stopwords. In more complex cases, it needs to **transform natural language queries into executable queries**. It may also accept user feedback and use it to **expand and refine** the original queries (relevance feedback).

- Indexer

It indexes the original raw documents in some data structures to enable efficient retrieval. The result is the **document index**. The most used data structure is called **inverted index** (easy to build and efficient to search).

- IR system

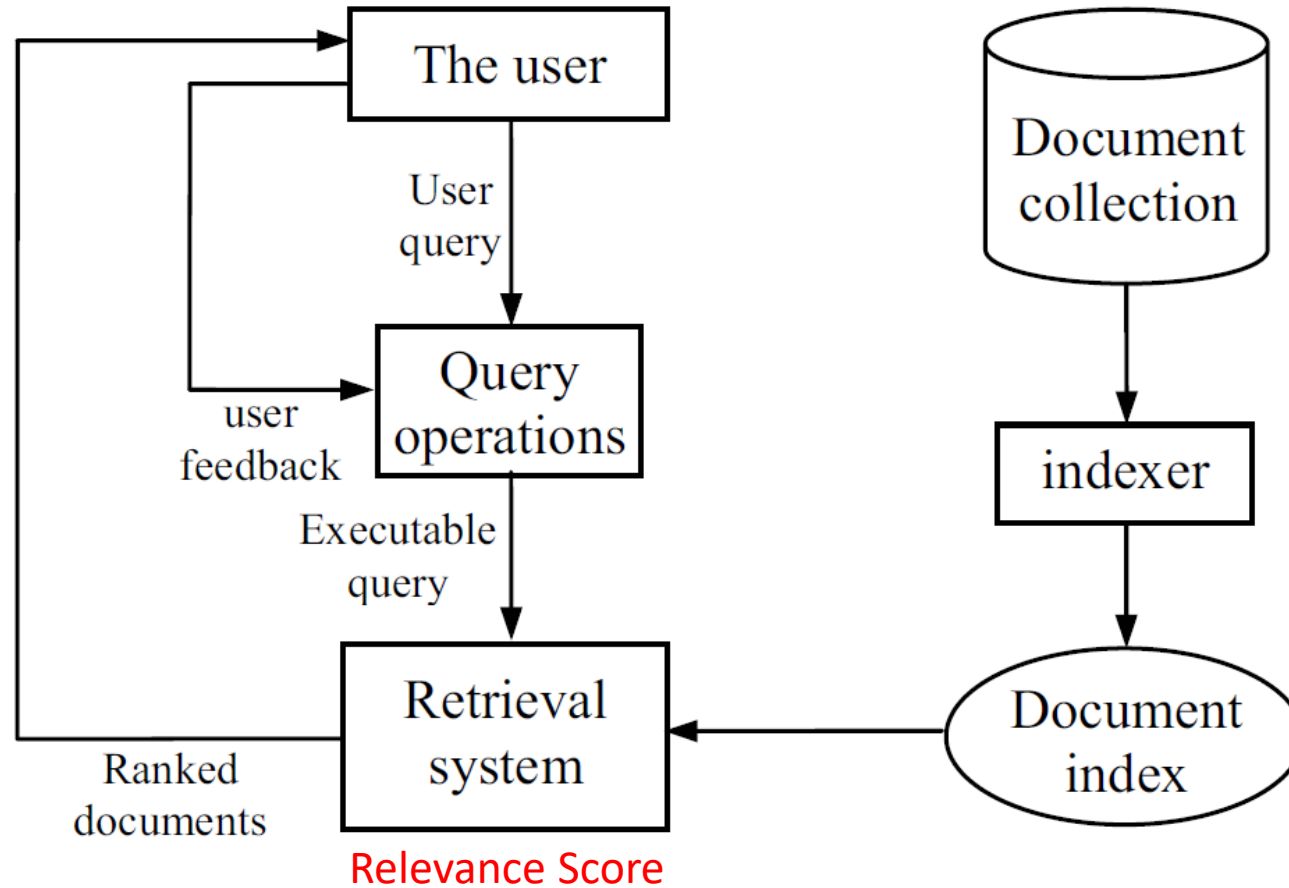
It computes a **relevance score for each indexed document to the query**. According to their relevance scores, the **documents are ranked and presented** to the user.

# Outline

- Goals and architecture of a IR system
- IR models
- Word2Vec and other modern Language Models



# Relevance Score



# IR models

- An IR model governs how a document and a query are represented and how the relevance of a document to a user query is defined.
- Main models:
  - Boolean model
  - Vector space model
  - Statistical language model
- They all treat each document or query as a “bag” of words or terms.
  - Term sequence and position in a sentence or a document are ignored.
  - Tokenization and other NLP preprocessing usually applies.

# Bag of words

- Given a collection of documents  $D$ , let  $V = \{t_1, t_2, \dots, t_{|V|}\}$  be the set of distinct words/terms in the collection.  $V$  is called the **vocabulary**.
- A **weight  $w_{ij} \geq 0$  is associated with each term  $t_i$  of a document  $\mathbf{d}_j \in D$ .** For a term that does not appear in document  $\mathbf{d}_j$ ,  $w_{ij} = 0$ .

$$\mathbf{d}_j = (w_{1j}, w_{2j}, \dots, w_{|V|j})$$

# Document-Term Matrix

- Bag-of-words models make possible to define the so-called **Document-Term Matrix**

[illegible]

# Boolean Model

- Document and query are represented by binary vectors whose length is  $|V|$  and their entries are such that

$$w_{ij} = \begin{cases} 1 & \text{if } t_i \text{ appears in } \mathbf{d}_j \\ 0 & \text{otherwise.} \end{cases}$$

- Query terms are combined logically using the Boolean operators **AND**, **OR**, and **NOT**.
  - E.g.,  $((data \text{ AND } mining) \text{ AND } (NOT \text{ text}))$
- **Exact match**: given a Boolean query, the system retrieves every document that makes the query logically true.
- The retrieval results are usually quite poor because term frequency is not considered.

# Retrieval in the "only-and" case (1/2)

## Documents

- D1 = "Hong Kong is in China"
- D2 = "Hong Kong is a beautiful city"
- D3 = "King Kong is a gorilla"

Preprocessing  
(lower-case + stop words removal + lemmatization)

## Vocabulary

hong, kong, china, beautiful, city,  
king, gorilla

## Document-Term Matrix

	hong	kong	china	beautiful	city	king	gorilla
D1	1	1	1	0	0	0	0
D2	1	1	0	1	1	0	0
D3	0	1	0	0	0	0	1

# Retrieval in the "only-and" case (2/2)

Document-Term Matrix  $M$

	hong	kong	china	beautiful	city	king	gorilla
D1	1	1	1	0	0	0	0
D2	1	1	0	1	1	0	0
D3	0	1	0	0	0	0	1

Identify Relevant Documents using Logic

D1 satisfies " $t_1$  and  $t_2$ "

D2 satisfies " $t_1$  and  $t_2$ "

D3 does not satisfy " $t_1$  and  $t_2$ "

**No rank in the produced results!!!**

Query

Hong Kong

Preprocessing + Vectorization  
(same as before)

$t_1$ =hong     $t_2$ =kong

# Vector Space Model

- Documents are also treated as a “bag” of words or terms.
- Each document is represented as a vector.
- However, the term weights are no longer 0 or 1. Each term weight is computed based on some variations of **TF** or **TF-IDF** scheme.
- **Term Frequency (TF) Scheme:** The weight of a term  $t_i$  in document  $\mathbf{d}_j$  is the number of times that  $t_i$  appears in  $\mathbf{d}_j$ , denoted by  $f_{ij}$ .
- Normalization may also be applied.



# TF: a very first example (using dot-product similarity, not cosine)

Document-Term Matrix  $M$

	hong	kong	china	beautiful	city	king	gorilla
D1	1	1	1	0	0	0	0
D2	1	1	0	1	1	0	0
D3	0	1	0	0	0	0	1

Query

Hong Kong

Preprocessing + Vectorization  
(same as before)

Vectorized Query  $q^T$

1	1	0	0	0	0	0
---	---	---	---	---	---	---

Calculate Relevant Documents

$$Mq =$$

Relevant  
Documents

2

2

1

Usual matrix-vector multiplication

# Vector Space Model

- The most well known weighting scheme

- TF: still **term frequency**

- IDF: **inverse document frequency**.

$N$ : total number of docs

$df_i$ : the number of docs where  $t_i$  appears.

- The final TF-IDF term weight is:

Normalized TF

$$tf_{ij} = \frac{f_{ij}}{\max\{f_{1j}, f_{2j}, \dots, f_{|V|j}\}}$$

$$idf_i = \log \frac{N}{df_i}$$

$$w_{ij} = tf_{ij} \times idf_i$$

# Retrieval in Vector Space Model

- Query  $\mathbf{q}$  is usually represented in the same way of documents.
- **Relevance of  $\mathbf{d}_j$  to  $\mathbf{q}$** : Compare the similarity of query  $\mathbf{q}$  and document  $\mathbf{d}_j$ .
- **Cosine similarity** (the cosine of the angle between the two vectors)

$$\text{cosine}(\mathbf{d}_j, \mathbf{q}) = \frac{\langle \mathbf{d}_j \bullet \mathbf{q} \rangle}{\|\mathbf{d}_j\| \times \|\mathbf{q}\|} = \frac{\sum_{i=1}^{|V|} w_{ij} \times w_{iq}}{\sqrt{\sum_{i=1}^{|V|} w_{ij}^2} \times \sqrt{\sum_{i=1}^{|V|} w_{iq}^2}}$$

- Another possibility is to adopt the **dot-product as similarity measure**
- On the contrary of Boolean model, the similarity allows to **rank the documents**

Note: Dot product can be used in place of cosine similarity, but the latter remains the preferred score function (recall also that when vectors are normalized, dot product is just a quick way to calculate cosine similarity)

# Example: Term Frequency Scheme

$d_1$	Text analysis is fun. I like doing text analysis.
$d_2$	I like doing text analysis. I like this. Puppies like this.
$d_3$	I like puppies, they are fun.
$d_4$	I like this blog post. This is fun.

	text	analysis	is	fun	I	like	puppies	doing	they	are	this	blog	post
$d_1$	2	2	1	1	1	1	0	1	0	0	0	0	0
$d_2$	1	1	0	0	2	3	1	1	0	0	2	0	0
$d_3$	0	0	0	1	1	1	1	0	1	1	0	0	0
$d_4$	0	0	1	1	1	1	0	0	0	0	2	1	1

$$d_1 = (2, 2, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0)$$

$$d_2 = (1, 1, 0, 0, 2, 3, 1, 1, 1, 0, 0, 2, 0, 0)$$

...

# Example: TF-IDF (1/3)

Raw Freq.		text	analysis	is	fun	I	like	puppies	doing	they	are	this	blog	post
$N=4$	$d_1$	2	2	1	1	1	1	0	1	0	0	0	0	0
	$d_2$	1	1	0	0	2	3	1	1	0	0	2	0	0
	$d_3$	0	0	0	1	1	1	1	0	1	1	0	0	0
	$d_4$	0	0	1	1	1	1	0	0	0	0	2	1	1
Normalized Term Freq.		text	analysis	is	fun	I	like	puppies	doing	they	are	this	blog	post
	$d_1$	1	1	0.5	0.5	0.5	0.5	0	0.5	0	0	0	0	0
	$d_2$	0.33	0.33	0	0	0.67	1	0.33	0.33	0	0	0.67	0	0
	$d_3$	0	0	0	1	1	1	1	0	1	1	0	0	0
	$d_4$	0	0	0.5	0.5	0.5	0.5	0	0	0	0	1	0.5	0.5
Document Frequency		text	analysis	is	fun	I	like	puppies	doing	they	are	this	blog	post
	$df_i$	2	2	2	3	4	4	2	2	1	1	2	1	1
TF-IDF		text	analysis	is	fun	I	like	puppies	doing	they	are	this	blog	post
	$d_1$	0.301	0.301	0.151	0.062	0.000	0.000	0.000	0.151	0.000	0.000	0.000	0.000	0.000
	$d_2$	0.099	0.099	0.000	0.000	0.000	0.000	0.099	0.099	0.000	0.000	0.202	0.000	0.000
	$d_3$	0.000	0.000	0.000	0.125	0.000	0.000	0.301	0.000	0.602	0.602	0.000	0.000	0.000
	$d_4$	0.000	0.000	0.151	0.062	0.000	0.000	0.000	0.000	0.000	0.000	0.301	0.301	0.301

$$w_{12} = \left( \frac{f_{12}}{\max\{f_{12}, f_{22}, \dots, f_{13,2}\}} \right) \times \log_{10} \frac{N}{df_1} = (0.33) \times \log_{10} \frac{4}{2} = 0.099$$

# Example: TF-IDF (2/3)

TF-IDF

	text	analysis	is	fun	I	like	puppies	doing	they	are	this	blog	post
$d_1$	0.301	0.301	0.151	0.062	0.000	0.000	0.000	0.151	0.000	0.000	0.000	0.000	0.000
$d_2$	0.099	0.099	0.000	0.000	0.000	0.000	0.099	0.099	0.000	0.000	0.202	0.000	0.000
$d_3$	0.000	0.000	0.000	0.125	0.000	0.000	0.301	0.000	0.602	0.602	0.000	0.000	0.000
$d_4$	0.000	0.000	0.151	0.062	0.000	0.000	0.000	0.000	0.000	0.000	0.301	0.301	0.301

Query,  $q = \{text, analysis\}$

	text	analysis	is	fun	I	like	puppies	doing	they	are	this	blog	post
$q$	0.301	0.301	0	0	0	0	0	0	0	0	0	0	0

$$w_{1q} = \left( \frac{f_{1q}}{\max\{f_{1q}, f_{2q}, \dots, f_{|V|q}\}} \right) \times \log_{10} \frac{N}{df_1} = \left( \frac{1}{1} \right) \times \log_{10} \frac{4}{2} = 0.301$$

$$w_{2q} = \left( \frac{f_{2q}}{\max\{f_{1q}, f_{2q}, \dots, f_{|V|q}\}} \right) \times \log_{10} \frac{N}{df_2} = \left( \frac{1}{1} \right) \times \log_{10} \frac{4}{2} = 0.301$$

# Example: TF-IDF (3/3)

TF-IDF

	text	analysis	is	fun	I	like	puppies	doing	they	are	this	blog	post
$d_1$	0.301	0.301	0.151	0.062	0.000	0.000	0.000	0.151	0.000	0.000	0.000	0.000	0.000
$d_2$	0.099	0.099	0.000	0.000	0.000	0.000	0.099	0.099	0.000	0.000	0.202	0.000	0.000
$d_3$	0.000	0.000	0.000	0.125	0.000	0.000	0.301	0.000	0.602	0.602	0.000	0.000	0.000
$d_4$	0.000	0.000	0.151	0.062	0.000	0.000	0.000	0.000	0.000	0.000	0.301	0.301	0.301

Query,  $q = \{text, analysis\}$

	text	analysis	is	fun	I	like	puppies	doing	they	are	this	blog	post
$q$	0.301	0.301	0	0	0	0	0	0	0	0	0	0	0

	$d_j \cdot q$	$\cos(d_j, q)$
$d_1$	0.181	0.886
$d_2$	0.060	0.495
$d_3$	0	0
$d_4$	0	0

# Latent Semantic Analysis (LSA)

- LSA = Applying **Truncated SVD** to the document-term matrix
- Motivations:
  - **Remove noise**
  - **Small storage** (dimension reduction)
  - **Recover missing data** via the help of co-occurrence patterns (addressed by the sparsity problem) = **synonym words share their weights to each other**



# Latent Semantic Analysis (LSA)

Index Words	Titles								
	T1	T2	T3	T4	T5	T6	T7	T8	T9
book			1	1					
dads						1			1
dummies		1						1	
estate							1		1
guide	1					1			
investing	1	1	1	1	1	1	1	1	1
market	1		1						
real							1		1
rich						2			1
stock	1		1					1	
value				1	1				

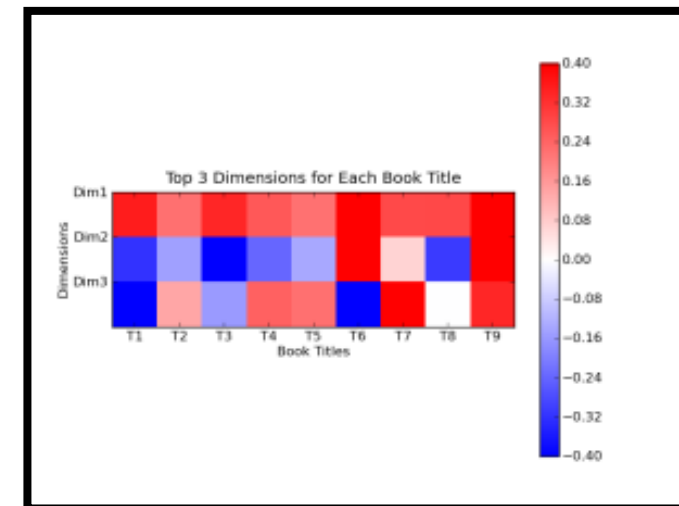
book	0.15	-0.27	0.04
dads	0.24	0.38	-0.09
dummies	0.13	-0.17	0.07
estate	0.18	0.19	0.45
guide	0.22	0.09	-0.46
investing	0.74	-0.21	0.21
market	0.18	-0.3	-0.28
real	0.18	0.19	0.45
rich	0.36	0.59	-0.34
stock	0.25	-0.42	-0.28
value	0.12	-0.14	0.23

 $*$ 

3.91	0	0
0	2.61	0
0	0	2

 $*$ 

T1	T2	T3	T4	T5	T6	T7	T8	T9
0.35	0.22	0.34	0.26	0.22	0.49	0.28	0.29	0.44
-0.32	-0.15	-0.46	-0.24	-0.14	0.55	0.07	-0.31	0.44
-0.41	0.14	-0.16	0.25	0.22	-0.51	0.55	0	0.34



# Latent Semantic Analysis (LSA)


- Document-term matrix  $X \in \mathbb{R}^{N \times |V|}$
- Extract **latent semantic** using Truncated SVD:

$$X \approx U_k \Sigma_k V_k^T$$

- Document-to-document matrix
  - $XX^T \approx U_k \Sigma_k V_k^T V_k \Sigma_k^T U_k^T = U_k \Sigma_k \Sigma_k^T U_k^T$
  - **document representation**:  $U_k \in \mathbb{R}^{m \times k}$
- Term-to-term matrix
  - $X^T X \approx V_k \Sigma_k^T U_k^T U_k \Sigma_k V_k^T = V_k \Sigma_k^T \Sigma_k V_k^T$
  - **term representation**:  $V_k \in \mathbb{R}^{n \times k}$

# Semantic Representations in LSA

- Rows of  $U_k$ , possibly multiplied by the singular values, are semantic representations of the terms/words
- Rows of  $V_k$  (i.e., columns of  $V_k^T$ ), possibly multiplied by the singular values, are semantic representations of the documents
- **Cosine similarity** can be calculated on these semantic representations in order to gauge the semantic similarity of terms or documents pairs
- **Clustering** algorithms can be executed on those representations in order to cluster together similar terms or documents
- Also **classification** algorithms can be trained and executed on those semantic representations



# Statistical Language Model

- Let:
  - the query  $q$  be a sequence of terms,  $q = q_1q_2...q_m$
  - the document collection  $D$  be a set of documents,  $D = \{d_1, d_2, ..., d_N\}$
- Consider the **probability of a query  $q$  as being “generated” by a probabilistic model based on a document  $d_j$ , i.e.,  $\Pr(q | d_j)$ .**
- Assuming **unigram model**, i.e., assume that each word is generated independently (practically, a multinomial distribution over words) so:

$$\Pr(q = q_1q_2...q_m | d_j) = \prod_{i=1}^m \Pr(q_i | d_j) = \prod_{i=1}^{|V|} \Pr(t_i | d_j)^{f_{iq}},$$

where  $f_{iq}$  is the number of times that term  $t_i$  occurs in  $q$ . Thus  $\Pr(t_i | d_j) = \frac{f_{ij}}{|d_j|}$ .

# Retrieval in Statistical Language Model

- To rank documents in retrieval, we are interested in estimating the posterior probability  $\Pr(d_j | q)$ . Using the **Bayes rule**, we have:

$$\Pr(d_j | q) = \frac{\Pr(q | d_j) \Pr(d_j)}{\Pr(q)}$$

- $\Pr(q)$  is not needed as it is the same for every document
- $\Pr(d_j)$  is usually considered uniform and thus will not affect ranking
- **We only need to compute the likelihood  $\Pr(q | d_j)$  as seen before!!!**

# Example of calculations

$$\prod_{i=1}^m P(q_i | d_j) = \prod_{i=1}^{|V|} P(t_i | d_j)^{f_{iq}}$$

$V = \{\text{"I"}, \text{"you"}, \text{"he"}, \text{"she"}, \text{"love"}, \text{"hate"}\}$

$q = \{\text{"I"}, \text{"love"}, \text{"I"}, \text{"love"}, \text{"you"}\}$

$$\begin{aligned} & \prod_{i=1}^m P(q_i | d_j) \\ &= P(\text{"I"} | d_j) \times P(\text{"love"} | d_j) \times P(\text{"I"} | d_j) \times P(\text{"love"} | d_j) \times P(\text{"you"} | d_j) \end{aligned}$$

$$\begin{aligned} & \prod_{i=1}^{|V|} P(t_i | d_j)^{f_{iq}} \\ &= P(\text{"I"} | d_j)^2 \times P(\text{"you"} | d_j)^1 \times P(\text{"he"} | d_j)^0 \times P(\text{"she"} | d_j)^0 \times \\ & \quad P(\text{"love"} | d_j)^2 \times P(\text{"hate"} | d_j)^0 \end{aligned}$$

# Smoothing in Statistical Language Models

- There is practical problem: a query term that does not appear in  $d_j$  has the probability of 0, which annihilate  $\Pr(q | d_j)$ .
- **Smoothing** is adopted to avoid this problem.

$$\Pr_{add}(t_i | d_j) = \frac{\lambda + f_{ij}}{\lambda |V| + |d_j|}.$$

where  $\lambda$  is a very low hyperparameter chosen in  $[0,1]$ .

- The name smoothing comes from the fact that these techniques tend to make distributions more uniform, by adjusting low probabilities such as zero probabilities upward, and high probabilities downward.

# Example

$$P(t_i|d_j) = \frac{f_{ij}}{|d_j|}$$

	text	analysis	is	fun	I	like	puppies	doing	they	are	this	blog	post
$d_1$	0.22	0.22	0.11	0.11	0.11	0.11	0.00	0.11	0.00	0.00	0.00	0.00	0.00
$d_2$	0.09	0.09	0.00	0.00	0.18	0.27	0.09	0.09	0.00	0.00	0.18	0.00	0.00
$d_3$	0.00	0.00	0.00	0.14	0.14	0.14	0.14	0.14	0.14	0.14	0.00	0.00	0.00
$d_4$	0.00	0.00	0.13	0.13	0.13	0.13	0.00	0.00	0.00	0.00	0.25	0.13	0.13

$$P_{add}(t_i|d_j) = \frac{1 + f_{ij}}{|V| + |d_j|}$$

	text	analysis	is	fun	I	like	puppies	doing	they	are	this	blog	post
$d_1$	0.14	0.14	0.09	0.09	0.09	0.09	0.05	0.09	0.05	0.05	0.05	0.05	0.05
$d_2$	0.08	0.08	0.04	0.04	0.13	0.17	0.08	0.08	0.04	0.04	0.13	0.04	0.04
$d_3$	0.05	0.05	0.05	0.10	0.10	0.10	0.10	0.10	0.10	0.10	0.05	0.05	0.05
$d_4$	0.05	0.05	0.10	0.10	0.10	0.10	0.05	0.05	0.05	0.05	0.14	0.10	0.10

Query,  $q$  = "like puppies post"

$$P(q|d_1) = 0.09 \times 0.05 \times 0.05 = 0.000225$$

$$P(q|d_2) = 0.17 \times 0.08 \times 0.04 = 0.000544$$

$$P(q|d_3) = 0.10 \times 0.10 \times 0.05 = 0.0005$$

$$P(q|d_4) = 0.10 \times 0.05 \times 0.10 = 0.0005$$

$$-\log P(q|d_1) = 3.648$$

$$-\log P(q|d_2) = \mathbf{3.264}$$

$$-\log P(q|d_3) = 3.301$$

$$-\log P(q|d_4) = 3.301$$

Minimize negative-log is just a trick that makes numbers "readable". Log-likelihood is also more efficient for calculations (sums in place of products)



# Limitation of the bag-of-words approaches

- Will the word ordering in the document affect the document representation?
- Yes!
- Considering *n*-grams allows to overcome this limitation ...
  - Example of bigrams (or 2-grams):

The text "Hong Kong is a beautiful city" contains the following bigrams:  
Hong\_Kong, Kong\_is, is\_a, a\_beautiful, beautiful\_city
- ... but it is computationally more expensive ... much more if  $n > 2$

# Outline

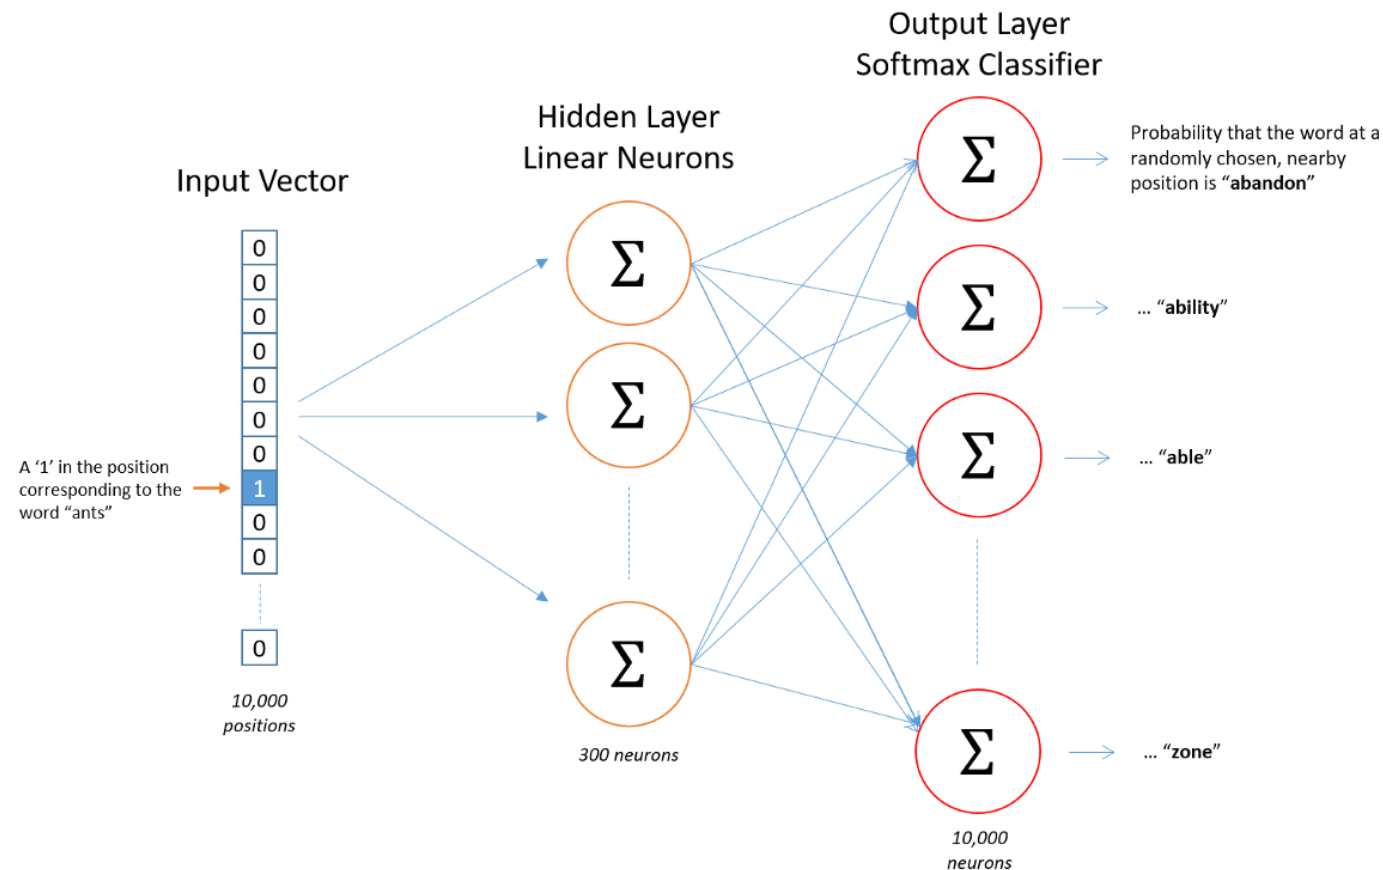
- Goals and architecture of a IR system
- IR models
- Word2Vec and other modern Language Models

# Vector representation used in Spacy?

- It uses **Word2Vec**, which generates a vector for every word
- **Word2vec explores the term co-occurrence in a local neighborhood (nearby words, and also called context words).**
- E.g., size of neighbourhood =2
  - “I like **this** **blog** **post**. **This** **is** fun.”.
  - Center word: brown (also called focus word).
  - Context words: blue (also called target words).
- Word2Vec considers all words as center words, and all their context words so that words can be predicted!
- **Linguistic principle: two words W and V are semantically related if they appear in similar contexts throughout a (very large) corpus of texts.**

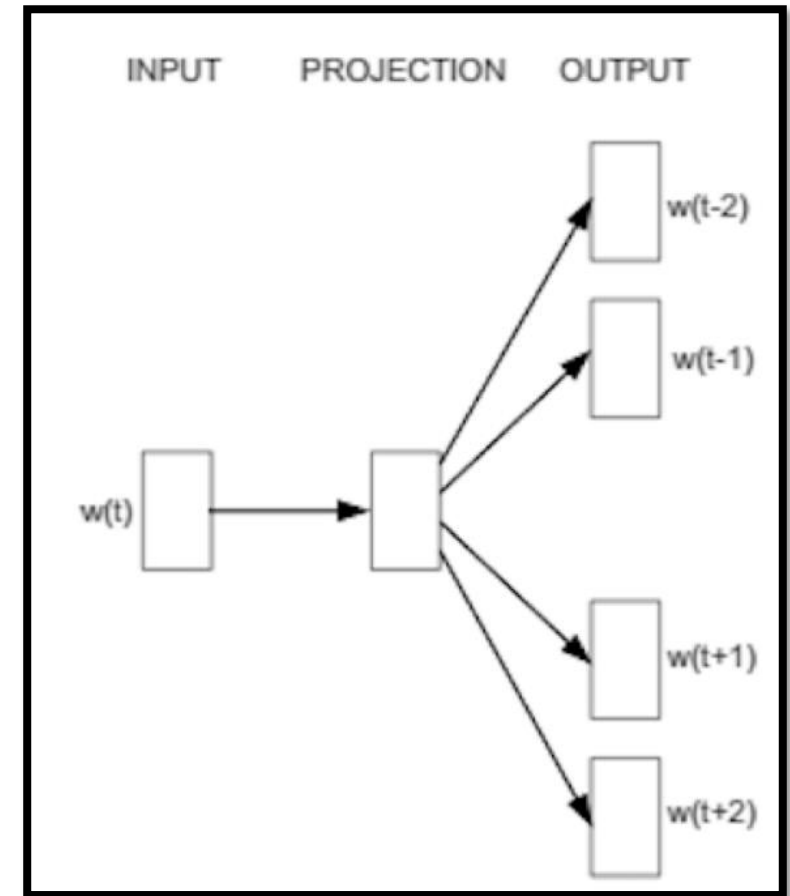
# Word2Vec (Skipgram)

- It uses an **AUTO-ENCODER-like neural network** trained on a large corpus




# Word2Vec (Skipgram)

- Maximize the **joint probability of predicting context words**  $w_{t-m}, \dots, w_{t-1}, w_{t+1}, \dots, w_{t+m}$  given  $w_t$
- $J(\theta) = \frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m, j \neq 0} \log p(w_{t+j} | w_t)$
- $p(w_{t+j} | w_t) = \frac{\exp(u_{w_{t+j}} \cdot v_{w_t})}{\sum_{w'} \exp(u_{w'} \cdot v_{w_t})}$
- Every word has 2 vectors
  - $v_w$  : when  $w$  is the center word
  - $u_w$  : when  $w$  is the context word
- Obtain the representation by maximization



# Word2Vec: Analogies emerge after learning



# From Word2Vec to Sentence Embeddings?

- Word vectors generated by Word2Vec can be aggregated to have a document vector
- But ... again word ordering does not affect the document vector
- Very recent Large Language Model based on neural networks allow to generate semantically meaningful document embedding by allowing the calculation of "contextual word embeddings" (same words has different embeddings based on the context where it appears) that can be "safely" aggregated to form semantically meaningful sentence embeddings
- Example:
  - SBERT =
    - = Sentence BERT =
    - = Sentence Bidirectional Encoder Representations from Transformers
    - <https://www.sbert.net/>

# References

- *Liu, Bing. Web data mining: exploring hyperlinks, contents, and usage data. Berlin: springer, 2011. Chapter 6.*
- *Tomas Mikolov, Ilya Sutskever, Kai Chen, Gregory S. Corrado, Jeffrey Dean. "Distributed Representations of Words and Phrases and their Compositionality," In Proceedings of NIPS 2013, pp. 3111-3119*
- <https://www.sbert.net/>